DS-GA 3001 007 │ **Lecture 2**

Reinforcement Learning

Jeremy Curuksu, PhD
NYU Center for Data Science
jeremy.cur@nyu.edu

February 3, 2023

DS-GA 3001 RL Curriculum

Reinforcement Learning:

- ▶ Introduction to Reinforcement Learning
- ▶ **Multi-armed Bandit**
- ▶ Dynamic Programming on Markov Decision Process
- ▶ Model free RL Prediction and Control
- ▶ Function Approximation and Deep RL
- ▶ Parametric Policy Function Optimization
- ▶ Planning from a RL Model
- ▶ Examples of Industrial Applications
- ▶ Advanced Topics and Development Platforms

Multi-armed Bandit

Last week:

- ▶ What is Reinforcement Learning?
- ▶ Key components of Reinforcement Learning
- ▶ Introduction to the Gym Python library

Today:

- ▶ **Multi-armed Bandit with action values**
- ▶ **Upper Confidence Bound**
- ▶ **Bayesian Bandit**
- ▶ **Policy Gradient Bandit**

Multi-armed Bandit with action values

What is Multi-armed Bandit?

What is Multi-armed Bandit?

The Multi-armed Bandit problem

- ▶ Reinforcement learning uses data it receives to evaluate actions (correct actions are not given), which creates a need to explore

- ▶ A Bandit is a RL problem involving learning to act in only one situation: 1 state, k possible actions

- ▶ No sequential structure, past actions do not influence the future: the distribution of reward r_t given a_t is identical and independent across time

What is Multi-armed Bandit?

Example of Multi-armed Bandit problem

	action	$ return
Week 1	A	
Week 2	B	
Week 3	?	

Which lever would you pull?

What is Multi-armed Bandit?

Example of Multi-armed Bandit problem

	action	$ return
Week 1	A	
Week 2	B	
Week 3	B	
Week 4	?	

How about now?

Exploration vs. Exploitation

Online decision-making involves a fundamental choice:

Exploitation:

Maximize performance using current knowledge

Exploration:

Increase knowledge

▶ The best strategy may involve short-term sacrifices

▶ The agent needs gather enough information to make the best overall decisions

Multi-armed Bandit Formalism

Problem Statement:

- ▶ The agent is faced repeateadly with a choice among k different actions ("*arms*")

- ▶ At each step t the agent selects an action a_t

- ▶ After each choice it receives a numerical reward r_t that depends on the action selected

- ▶ The distribution $p(r|a)$ is fixed but unknown

- ▶ Goal is to maximize cumulative reward:

$$\sum_{i=1}^{t} r_i$$

Exploit knowledge with action value

Action value for action a is the expected reward:

$$q(a) \doteq \mathbb{E}(r|a) = p(r|a) \times r = \lim_{t \to +\infty} \frac{1}{t} \sum_{i=1}^{t} r_i | a$$

▶ An estimate is the average of the sampled rewards:

$$q_t(a) \doteq \frac{\textit{sum of rewards when a taken prior to t}}{\textit{number of times a taken prior to t}}$$

▶ With an estimate of $q(a)$, we can select an action:

Greedy policy: $\quad a_t \doteq \arg\max_a q_t(a)$

Incremental implementation

The agent can learn online with a moving average:

For $a = a_t$:
$$q_t(a) = \frac{1}{t} \sum_{i=1}^{t} r_i | a$$

$$q_t(a) = q_{t-1}(a) + \frac{1}{t} (r_t - q_{t-1}(a))$$

$\forall \, a \neq a_t$:
$$q_t(a) = q_{t-1}(a)$$

For non-stationary problems, the agent can *track* q(a):

$$q_t(a) = q_{t-1}(a) + \alpha (r_t - q_{t-1}(a))$$

Explore new actions with ϵ-greedy

The agent must explore to learn q-values

- ► Greedy selection always exploits current knowledge on q-values to maximize reward, it never explore

- ► Alternative: Behave greedily most of the time, but every once in a while select a random action

- ► **ϵ-greedy algorithm:**
    - ► Select random action (explore) with $p = \epsilon$
    - ► Select greedy action (exploit) with $p = 1 - \epsilon$

- ► When all actions are sampled indefinitely:

$$\lim_{t \to +\infty} q_t(a) = q(a)$$

Practice: k-armed Bandit Algorithm

k-armed Bandit both evaluates $q(a)$ and improves a:

```
Initialize, for a = 1 to k:
  q(a) = 0
  n(a) = 0

Loop forever:
  a  = random action  with p = epsilon
  or = argmax q(a)     with p = 1 - epsilon
  Execute a, observe r
  n(a) = n(a) + 1
  q(a) = q(a) + 1/n(a) * (r - q(a))
```

Case Study: 10-armed testbed

k-armed Bandit problem
Distribution of $q(a)$ vs. action a

Average performance of ϵ-greedy
Average reward over 2000 runs vs. time step

* Sutton and Barto, 1998

Total regret L_t

Analyzing regret in Multi-armed Bandit

▶ How can we reason about the exploration trade off?

▶ The (true) optimal value is: $v_* = \max_a q(a)$

▶ Regret is the opportunity loss at step t: $v_* - q(a_t)$

▶ Thus the best trade-off between exploration and exploitation is the one that minimizes total regret L_t:

$$L_t = \sum_{i=1}^{t} (v_* - q(a_i))$$

▶ The agent cannot measure regret directly, but regret can be used to analyze different RL algorithms on solved problems

Action Regret Δ_a

Analyzing regret in Multi-armed Bandit

▶ The action regret Δ_a for an action a is the difference between the optimal value and the true value of a:

$$\Delta_a = (v_* - q(a))$$

▶ Total regret can be defined by action regrets and action counts:

$$L_t = \sum_{i=1}^{t} (v_* - q(a_i)) = \sum_{a \in (A)} N_t(a)(v_* - q(a)) = \sum_{a \in (A)} N_t(a)\Delta_a$$

▶ Thus the best trade-off between exploration and exploitation is the one that ensures small count for actions with large regret

Upper Confidence Bound

Explore new actions with UCB

Upper Confidence Bound (UCB)

▶ For each action value $q(a)$, compute an upper confidence $u_t(a)$ such that $q(a) \leq q_t(a) + u_t(a)$

▶ Select action that maximizes this Upper Confidence Bound:

$$a_t = \arg\max_{a \in (A)} [q_t(a) + u_t(a)]$$

where:

$$u_t(a) = c \sqrt{\frac{\ln t}{N_t(a)}}$$

▶ The UCB algorithm can achieve logarithmic expected total regret (demonstration out of scope)

Explore new actions with UCB

Upper Confidence Bound (UCB)

$$a_t = \underset{a \in (A)}{\arg\max} \left[q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

▶ Uncertainty depends on number of times an action is selected:
  ▶ **Small $N_t(a)$ ⇒ Large $u_t(a)$** ⇐ Estimated q-value uncertain
  ▶ **Large $N_t(a)$ ⇒ Small $u_t(a)$** ⇐ Estimated q-value is accurate

▶ UCB favors an action because its estimated q-value is high, or because it has not been explored a lot relative to time elapsed

▶ UCB guarantees all actions will be explored without the need to manually predefine an ϵ-schedule

Case Study: 10-armed testbed

Average performance of ϵ-greedy and UCB algorithms

Average reward over 2000 runs vs. time step

Bayesian Bandit

The Bandit Model

A Bandit model is a reward transition function:

$$p(r \mid a) = p(r_{t+1} = r \mid a_t = a) \iff r(a) = \mathbb{E}(r, a)$$

where $\quad \mathbb{E}(r|a) = p(r|a) \times r = \lim_{t \to +\infty} \sum_{i=1}^{t} \hat{r}_i = \lim_{t \to +\infty} \frac{1}{t} \sum_{i=1}^{t} r_i$

The Bandit Model

A Bandit model is a reward transition function:

$$p\left(r \mid a\right) = p\left(r_{t+1} = r \mid a_t = a\right) \;\Leftrightarrow\; r(a) = \mathbb{E}(r, a)$$

where $\quad \mathbb{E}(r|a) = p(r|a) \times r = \lim_{t \to +\infty} \sum_{i=1}^{t} \hat{r}_i = \lim_{t \to +\infty} \frac{1}{t} \sum_{i=1}^{t} r_i$

▶ **Bandit model-based algorithm:** (Expectation model)

$$\hat{r}_t(a) = \hat{r}_{t-1}(a) + \alpha \left(r_t - \hat{r}_{t-1}(a)\right)$$

The Bandit Model

A Bandit model is a reward transition function:

$$p\left(r \mid a\right) = p\left(r_{t+1} = r \mid a_t = a\right) \;\Leftrightarrow\; r(a) = \mathbb{E}(r, a)$$

where $\quad \mathbb{E}(r|a) = p(r|a) \times r = \lim_{t \to +\infty} \sum_{i=1}^{t} \hat{r}_i = \lim_{t \to +\infty} \frac{1}{t} \sum_{i=1}^{t} r_i$

▶ **Bandit model-based algorithm:** (Expectation model)

$$\hat{r}_t(a) = \hat{r}_{t-1}(a) + \alpha \left(r_t - \hat{r}_{t-1}(a)\right)$$

▶ **Bandit value-based algorithm:**

$$q_t(a) = q_{t-1}(a) + \alpha \left(r_t - q_{t-1}(a)\right) \quad \text{...Identical?}$$

Bayesian Bandit Models

**Bayesian Bandit models the full distribution of rewards:**

- ▶ Bayesian Bandit tracks a parameterized distribution function of expected reward $p(\mathbb{E}(r)|\theta, a)$, called likelihood function

- ▶ Selects actions based on $p(\mathbb{E}(r)|\theta, a)$ e.g., using UCB

- ▶ Uses reward observed to update posterior distributions of θ:

$$p_t(\theta|r) \propto p(\mathbb{E}(r)|\theta, a) \times p_{t-1}(\theta|r)$$

- ▶ For example, $\theta = (\mu, \sigma)_a$ if $p(\mathbb{E}(r)|\theta, a)$ are Gaussian distributions

- ▶ Bayesian Bandit allows to inject prior knowledge $p_0(\theta|a)$ and then use posteriors to guide exploration

Example: Bayesian Bandit with UCB

Apply UCB to a Bayesian Bandit model:

- ▶ Define likelihood function $p(\mathbb{E}(r)|\theta, a) = p_\theta(q(a))$ as Gaussians
- ▶ Estimate upper confidence from posterior using $u_t(a) = c\sigma_t(a)$ where $\sigma_t(a)$ is the standard deviation of $p_\theta(q(a))$
- ▶ Select action $a_t = \arg\max_a \left(q_t(a) + c\sigma_t(a) \right)$

Bandit with Thompson Sampling

Bayesian model with Probability Matching:

- ▶ Instead of selecting actions from q-values with highest mean according to $p_\theta(q(a))$ with ϵ-greedy or UCB, Thompson sampling explicitly samples q-values from $p_\theta(q(a))$

- ▶ Thompson sampling selects action a according to probability that a is the optimal action given the data sampled so far

$$\pi_t(a) = p\left(q(a) = \max_{a'} q(a') \mid \text{history}_{t-1}\right)$$

$$\pi_t(a) = \mathbb{E}\left(\mathcal{I}\left(q_t(a) = \max_{a'} q_t(a')\right) \mid \text{history}_{t-1}\right)$$

$$\pi_t(a) = \sum_{i=1}^{t} \frac{1}{t}\left(\mathcal{I}\left(q_t(a) = \max_{a'} q_t(a')\right)\right)$$

where $\quad \mathcal{I}(\text{True}) = 1, \quad \mathcal{I}(\text{False}) = 0$

Toward Sequential RL and MDP...

Information State Space Bandit Model

- ▶ Bayesian Bandit tracks an evolving probability distribution of reward, which can be considered an *information state s_t*

- ▶ Each action a_t causes a transition to a new state s_{t+1} (by adding information), which is a sequential RL problem

- ▶ The tree of possible chains of events grows extremely rapidly, so approximate RL methods (lecture 5) are required

Contextual Bandits

- ▶ If context on distinctive states are given to the agent, it can learn actions and values specific to each state

- ▶ Simplified case of more general sequential RL problem where actions may affect next states and thus future possible rewards

Policy Gradient in Multi-armed Bandit

Policy Gradient Bandit

Can we learn a policy without learning values?

► Yes we can!

► Define a parameterized function $\pi_\theta(a) : a \mapsto p_\theta(a)$ and learn parameters θ that maximize a performance measure $J_{\pi_\theta}(a)$

► $\pi_\theta(a)$ can be arbitrary (just need distinguish possible actions)

► $J_{\pi_\theta}(a)$ can also be arbitrary (e.g. "always turn right in a maze")

► If $J_{\pi_\theta}(a)$ is unknown, it needs to be learned... It is often defined based on $q_t(a) =$ the *critic* in Actor-Critic algorithms:

$$\theta = \theta + \alpha \nabla_\theta q(a)$$

► **Out of scope for today (covered in depth in lecture 6)**

Thank you!