

# DS-GA 3001 005 | Lecture 3

## Reinforcement Learning

---

Jeremy Curuksu, PhD

NYU Center for Data Science

[jeremy.cur@nyu.edu](mailto:jeremy.cur@nyu.edu)

February 14, 2024

# DS-GA 3001 RL Curriculum

---

## Reinforcement Learning:

- ▶ Introduction to Reinforcement Learning
- ▶ Multi-armed Bandit
- ▶ **Dynamic Programming on Markov Decision Process**
- ▶ Model-free Reinforcement Learning
- ▶ Value Function Approximation (Deep RL)
- ▶ Policy Function Approximation (Actor-Critic)
- ▶ Planning from a Model of the Environment
- ▶ Examples of Industrial Applications
- ▶ Advanced Topics and Development Platforms

# Dynamic Programming on Markov Decision Process

---

## Last week:

- ▶ Multi-armed Bandit with action values ( $\epsilon$ -greedy)
- ▶ Upper Confidence Bound
- ▶ Bayesian Bandit

## Today:

- ▶ **Markov Decision Process**
- ▶ **Value Functions and Bellman Equations**
- ▶ **Dynamic Programming**

# Generalization to Sequential RL


---

## Sequential Goal-Directed Reinforcement Learning

- ▶ Bandit problems have only one state, but often the agent must learn different actions in different situations (states)
- ▶ Actions in turn may influence subsequent states, and through those states may influence future rewards
- ▶ To learn to make good decisions, we need assign credit for long term consequences to individual actions

# Markov Decision Process

# Markov Decision Process (MDP)



## At each step, the agent:

- Finds itself in state  $s_t$  (from  $o_t$ )
- Executes action  $a_t$
- Receives reward  $r_{t+1}$

## The environment:

- Receives action  $a_t$
- Send reward  $r_{t+1}$
- Send observation  $o_{t+1}$

# Markov Decision Process (MDP)

An MDP is a mathematical idealization of goal-directed learning from interaction with an environment

- ▶ Simulating a MDP produces a sequence of  $n$  tuples (trajectory)

$$(s_t, a_t, r_{t+1}, s_{t+1})_n = (s_0, a_0, r_1, s_1, a_1, r_2, \dots, s_n)$$

- ▶ The environment dynamics is fully characterized by the joint probability of each possible  $s_{t+1}$  and  $r_{t+1}$  as a function of the immediately preceding state and action,  $s_t$  and  $a_t$

$$p(s', r | s, a) = p(s_{t+1} = s', r_{t+1} = r | s_t = s, a_t = a)$$

- ▶ **Markov property:** *The state must include all information from past agent–environment interactions that influence the future*

$$p(s, r | s_t, a_t) = p(s, r | H_t, a_t)$$

# Goals and Rewards

## RL applies the reward hypothesis

- ▶ The purpose of an RL agent is formalized in term of a signal called *reward*  $r_t \in \mathbb{R}$  passing from the environment to the agent
- ▶ The agent goal is to maximize the amount of reward it receives

## Reward:

$$r_t$$

## Optimal Policy:

$$\pi_* = \arg \max_a \left( \sum r_t \right)$$



# Agent goal is to maximize return $G_t$

$G_t$  is the total accumulated reward from time-step  $t$

- ▶ Acting in a MDP results in returns  $G_t$  that depend on the policy:

$$G_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots + r_T$$

- ▶  $G_t$  can be discounted by factor  $\gamma \in [0, 1]$  to account for present value of future rewards (in episodic or continuing tasks)

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+1+k}$$

- ▶  $\gamma < 1 \Rightarrow$  Immediate rewards  $>$  delayed rewards
- ▶  $\gamma$  close to 0  $\Rightarrow$  "Myopic" agent
- ▶  $\gamma$  close to 1  $\Rightarrow$  "Far-sighted" agent

# **Value Functions and Bellman Equations**

# State Value Function $v_{\pi}(s)$

## Expected return when starting in $s$ and following $\pi$

- Rewards the agent can expect to receive in the future depend on what actions it will take. Accordingly, value functions are defined with respect to particular ways of acting (policies)

$$\forall s \in \mathcal{S}, \quad v_{\pi}(s) \doteq \mathbb{E}_{\pi}(G_t | s)$$

$$v_{\pi}(s) = \mathbb{E}_{\pi}(r_{t+1} + \gamma G_{t+1} | s)$$

$$v_{\pi}(s) = \mathbb{E}(r_{t+1} + \gamma v_{\pi}(s_{t+1}) | s)$$

$$v_{\pi}(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_{\pi}(s')]$$

- $v_{\pi}(s)$  indicates how good it is to be in  $s$  when following  $\pi$

# Action Value Function $q_{\pi}(s, a)$

## Expected return when selecting $a$ in $s$ and following $\pi$

- The action value more directly informs on which action to take

$$\forall s \in \mathcal{S}, \quad q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}(G_t | s, a)$$

$$q_{\pi}(s, a) = \mathbb{E}(r_{t+1} + \gamma q_{\pi}(s_{t+1}, a_{t+1}) | s, a)$$

$$q_{\pi}(s, a) = \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \sum_{a'} \pi(a' | s') q_{\pi}(s', a') \right]$$

- $q_{\pi}(s, a)$  indicates how good it is to select  $a$  in  $s$  under  $\pi$
- Note that  $\sum_a \pi(a | s) q_{\pi}(s, a) = \mathbb{E}(q_{\pi}(s, a)) = v_{\pi}(s) \quad \forall s$

# Optimal Value Functions $v_*$ and $q_*$

## Bellman Optimality equations

- ▶  $v_*(s)$  is the maximum state-value function over all policies:

$$v_*(s) \doteq \max_{\pi} v_{\pi}(s)$$

$$v_*(s) = \max_a \mathbb{E}(r_{t+1} + \gamma v_*(s_{t+1}) \mid s, a)$$

$$v_*(s) = \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')]$$

- ▶  $q_*(s, a)$  is the maximum action-value function over all policies:

$$q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$$

$$q_*(s, a) = \mathbb{E}(r_{t+1} + \gamma \max_{a'} q_*(s_{t+1}, a') \mid s, a)$$

$$q_*(s, a) = \sum_{s', r} p(s', r \mid s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]$$

# Summary of Bellman equations

There are four main Bellman equations:

$$v_{\pi}(s) = \mathbb{E}(r_{t+1} + \gamma v_{\pi}(s_{t+1}) \mid s) \quad (1)$$

$$v_{*}(s) = \max_a \mathbb{E}(r_{t+1} + \gamma v_{*}(s_{t+1}) \mid s, a) \quad (2)$$

$$q_{\pi}(s, a) = \mathbb{E}(r_{t+1} + \gamma q_{\pi}(s_{t+1}, a_{t+1}) \mid s, a) \quad (3)$$

$$q_{*}(s, a) = \mathbb{E}(r_{t+1} + \gamma \max_{a'} q_{*}(s_{t+1}, a') \mid s, a) \quad (4)$$

► There are equivalences between state and action values:

$$v_{\pi}(s) = \sum_a \pi(a \mid s) q_{\pi}(s, a) = \mathbb{E}(q_{\pi}(s, a))$$

$$v_{*}(s) = \max_{\pi} v_{\pi}(s) = \max_a q_{*}(s, a)$$

# Policy Evaluation and Optimization

## Bellman equations are used for prediction and control

- **Prediction:** Evaluate a policy by estimating  $v_\pi$  or  $q_\pi$

$$v_\pi(s) \geq v_{\pi'}(s) \iff \pi \geq \pi' \quad \forall s$$

- **Control:** Improve a policy based on  $v_\pi$  or  $q_\pi$

- **Example:**  $\pi_*(s, a) = \begin{cases} 1, & \text{if } a = \arg \max_a (q_*(s, a)) \\ 0, & \text{otherwise} \end{cases}$

- **Theorem:** For any MDP, there exists an optimal policy  $\pi_*$  that is better than or equal to all other policies:  $\pi_* \geq \pi, \forall \pi$
- There is always at least one deterministic optimal policy for any MDP. There can be multiple optimal policies

# Solving Bellman Equations

## Solving the RL Prediction problem

- ▶ Bellman equations are linear so can in principle be solved:

$$V = R + \gamma P^\pi V$$

$$(I - \gamma P^\pi) V = R$$

$$V = (I - \gamma P^\pi)^{-1} R$$

where:  $v_i = v(s_i)$ ,  $r_i = \mathbb{E}_\pi[r_t | s_i]$ ,  $P_{ij}^\pi = \sum_a \pi(a | s_i) p(s_j | s_i, a)$

- ▶ Solving Bellman equations algebraically is akin to exhaustive search ( $O(|s|^3)$ ), it can be computed only for small problems
- ▶ This method assumes (1) Markov property, (2) MDP dynamics is known, (3) we have enough resources to compute the solution



# Solving Bellman Equations

---

## Solving the RL Prediction problem

- ▶ Bellman equations are linear so can in principle be solved:

$$V = (I - \gamma P^\pi)^{-1} R$$

## Solving the RL Optimization problem

- ▶ Bellman optimality equations are non-linear thus can't be solved directly
- ▶ RL optimization relies on iterative solution methods
  - ▶ Dynamic Programming (use a model)
  - ▶ Monte-Carlo, Temporal Difference (use samples)

# Dynamic Programming

# Dynamic Programming

---

- ▶ DP refers to a collection of algorithms to evaluate and/or improve policies given a model of the environment as a MDP
- ▶ DP is an essential foundation: all RL methods can be viewed as attempts to achieve the same effect as DP, but with less computation and without a perfect model of the environment
- ▶ Key idea of DP is the use of value functions to organize the search for good policies
- ▶ All DP methods consist of two different parts: policy evaluation and (optionally) policy improvement
- ▶ All DP methods update estimates of the values of states based on estimates of the values of successor states (*bootstrapping*)

# Policy Evaluation for a Given Policy

## Estimate $v_\pi(s)$ of a given policy $\pi$

- Turn the Bellman equation

$$v_\pi(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$

...into an update function:

- **Initialize  $v_0$  e.g., to zero, then iterate:**

$$\forall s, v_{k+1}(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')]$$

- **Whenever  $v_{k+1}(s) = v_k(s)$ , for all  $s$ , we have found  $v_\pi$**
- It can be shown that  $\lim_{k \rightarrow \infty} v_k = v_\pi$  (demonstration out of scope)

# Example of Policy Evaluation



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

$R_t = -1$   
on all transitions

## Evaluate random policy $\pi_{\text{random}}$ in this $4 \times 4$ gridworld

- At each iteration  $k$ , loop through all states and update value estimate  $v_k(s)$  of every state  $s$  for  $\pi_{\text{random}}$

# Example of Policy Evaluation

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0


$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

# Example of Policy Evaluation



# Policy Improvement

**Find a better policy  $\pi'$  given  $v_\pi(s)$**

**1. For a given policy  $\pi$ , compute:**

$$\forall s : \pi'(s) = \arg \max_a q_\pi(s, a) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$

**2. Evaluate  $v_{\pi'}(s)$  as in previous slides (policy evaluation)**

**3. Repeat**

**Policy Improvement Theorem:**

$$\forall s, v_{\pi'}(s) = \max_a q_\pi(s, a) \geq v_\pi(s) \implies \pi' \text{ better or same as } \pi$$

- ▶ When  $v_{\pi'}(s) = v_\pi(s)$ ,  $v_{\pi'} = \max_a q_{\pi'}(s, a)$ . This is the Bellman optimality equality, thus  $\pi'$  is optimal.
- ▶ Thus, if  $v_{\pi'}(s) \geq v_\pi(s)$ ,  $\pi'$  either is an improvement or is optimal




# Generalized Policy Iteration

All RL methods are Generalized Policy Iteration methods

Control

Improve Policy



# Practice: Policy Iteration Algorithm

Policy Iteration iterates multiple loops over all states to evaluate  $v$ , then loops over all states once to improve  $\pi$ , then repeats:

Initialize  $v(s)$  and  $\pi(s)$  arbitrarily for all  $s$

1. Loop:

$$\Delta = 0$$

For each  $s$ :

$$v_{\text{old}} = v(s)$$

$$v(s) = \sum_a \pi(a | s) \sum_{s', r} p(s', r | s, a) [(r + \gamma v(s'))]$$

$$\Delta = \max(\Delta, |v_{\text{old}} - v(s)|)$$

Stop when  $\Delta < \xi$

2. For each  $s$ :

$$\pi_{\text{old}}(s) = \pi(s)$$

$$\pi(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [(r + \gamma v(s'))]$$

Stop if  $\pi_{\text{old}} \iff \pi(s)$ , else go to step 1

# Practice: Value Iteration Algorithm

## Policy improvement with truncated policy evaluation

- ▶ Policy iteration involves policy evaluation at each iteration, which may itself require multiple loops through all states
- ▶ Is exact convergence needed, or can we stop sooner? When?
- ▶ Policy evaluation can be truncated in several ways without losing the convergence guarantees of policy iteration
- ▶ **A special case is when policy evaluation is stopped after just one loop (one update of each state). It is equivalent to turning the Bellman optimality equation into an update function:**

$$v_{k+1}(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_k(s')]$$

# Practice: Value Iteration Algorithm

Value Iteration truncates policy evaluation to 1 step between two (greedy) policy improvement steps while looping over all states

Initialize  $v(s)$  arbitrarily for all  $s$

Loop:

$$\Delta = 0$$

For each  $s$ :

$$v_{\text{old}} = v(s)$$


$$v(s) = \max_a \sum_{s',r} p(s',r | s, a) [(r + \gamma v(s'))]$$

$$\Delta = \max(\Delta, |v_{\text{old}} - v(s)|)$$

Stop when  $\Delta < \xi$

$$\pi(s) = \arg \max_a \sum_{s',r} p(s',r | s, a) [(r + \gamma v(s'))]$$

# Example of Value Iteration



# Asynchronous Dynamic Programming

## Update values in any order whatsoever...

- ▶ DP algorithms described so far loop over all states, but in practice this is often impossible (e.g., Chess has  $10^{40}$  states)
- ▶ Asynchronous DP backs up states in any order, and still converges if it continues to update values of all states
- ▶ **Asynchronous DP makes it possible to focus DP updates onto parts of the state space that are most relevant to the agent:**
  - ▶ **Prioritised Sweeping:** States with largest Bellman Error:

$$\max |[r_{t+1} + \gamma \hat{v}(s_{t+1}) | s] - \hat{v}(s)|$$

- ▶ **Real-time DP:** Agent's real experience determines states to update, while latest values guide its decision making


# Efficiency of Dynamic Programming

---

**DP provides a well defined notion of optimality, but is often an ideal that AI agents can only approximate**

- ✓ Asynchronous DP often exponentially faster than direct search
- ✓ ...in particular if agent starts with good initial values or policies
- ✓ DP is iterative so can learn with limited compute resources
- ✓ With today's computers, DP can solve MDPs with millions of state (assuming a small number of actions)
- ✗ In most cases of practical interest, a perfect MDP model of state transitions and rewards is not available
- ✗ In most cases of practical interest, there are far more states than there could possibly be entries in a look-up table

# Efficiency of Dynamic Programming



## DP often suffers from the curse of dimensionality

- ▶ DP uses full-width backups
- ▶ Even one full-depth backup can be too expensive
- ▶ Need to sample (next lecture)



# Today's Takeaways

---

**Dynamic Programming provides a foundation for RL by numerically solving the Bellman equations through GPI, which involves two processes taking place in parallel:**

- ▶ **Process of prediction** which can learn values online at every step by bootstrapping estimates on the basis of previous estimates until they are consistent with a policy followed.
- ▶ **Process of optimization** which can improve the policy followed by making it greedy with respect to the latest value estimates.
- ▶ **Both processes stabilize only when a policy has been found that is greedy with respect to its own value function.**
- ▶ But Dynamic Programming needs a model of the environment.

**Thank you!**